

# Introduction to PRAM

Oguz Kaya

Assistant Professor  
Université Paris-Saclay and ParSys team at LISN, Orsay, France

# Outline

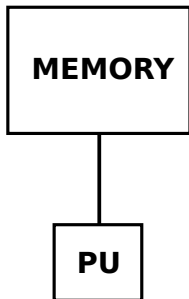
## 1 Introduction to PRAM

## 1 Introduction to PRAM

- université  
PARIS-SACLAY

# RAM model

RAM (random access machine) model is an abstraction of computers



- Consists of a processing unit (PU) and an associated memory space.
- Access to each element in the memory is done in constant time.
- Each operation is performed in constant time in the PU.
- Equivalent to a turing machine (simulation)
- Belongs to the class “register machines”
- Not very realiste (L1/L2/L3/DRAM memory hierarchy) yet useful for developing algorithms that are “asymptotically optimal”.
  - One should then adapt these algorithms to modern computer architectures, i.e. HPC.

Algorithm: A sequence of operations on a RAM machine, which aims to solve a given problem.

```


1:  $sum = 0$ 
2: for  $i = 1, \dots, N$  do
3:    $sum \leftarrow sum + A[i]$ 
4: return  $sum$ 

```

- Each memory access and arithmetic operation takes unit time.
- The **time complexity** of the algorithm is a function  $T(N)$  of the problem size which represents the total number of operations it performs.
  - One can also consider its **space complexity** which represents the total memory space it uses in execution.

## Complexity analysis

Many methods exist for complexity analysis:

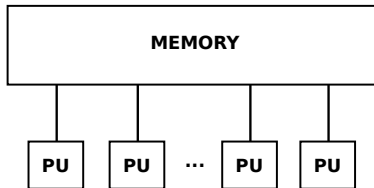


- Nested loops: Counting techniques in the loop iteration domain
- Recursion / divide-and-conquer: Counting by hand, Master theorem
- Probabilistic analysis (i.e., quicksort)
- Amortized analysis (i.e., dynamic tables)
- These analysis could be performed for
  - the best case
  - the worst case
  - the average case

# PRAM model

Parallel random access machine (PRAM) is an abstraction of parallel computers.


- Consists of  $P$  PUs and a memory space shared among all PUs.



- PUs
  - execute the same algorithm/code in a synchronous manner, but on different elements in the memory (i.e., à la SIMD).
  - can perform an arithmetic operation **simultaneously** in constant time.
  - can do reads/writes from/to different memory locations **simultaneously** in constant time.




Parallel random access machine (PRAM) is an abstraction of parallel computers.



- Read/write on the same memory location?
  - This behavior defines different PRAM models.
- Number of PUs: Potentially infinite!
- Memory size: Potentially infinite!
- Communication cost: Completely ignored (done through fast shared memory implicitly)
  - PRAM is useful for developing “asymptotically optimal” parallel algorithms
  - One should then adapt these algorithms to modern parallel computer architectures.

# Memory conflicts in PRAM

During the execution of a PRAM algorithm, multiple PUs could try to access to the same memory location (memory conflict).




- Possible resolutions?
- **CREW:** Concurrent Read Exclusive Write
  - Simultaneous reading of the same memory location is allowed, and takes constant time
  - Simultaneous writing is forbidden (bug)
  - Standard PRAM model, closest to real machines

# Memory conflicts in PRAM


During the execution of a PRAM algorithm, multiple PUs could try to access to the same memory location (memory conflict).

- Possible resolutions?
- **CRCW:** Concurrent Read Concurrent Write
  - Simultaneous reading/writing from/to the same memory location is allowed, and takes constant time
  - The most powerful model
  - A “concurrency mode” is employed for simultaneous writes:
    - Consistent mode: All PUs write the same value.
    - Arbitrary mode: One PUs value is retained randomly.
    - Priority mode: The PU with min/max index's value is retained.
    - Fusion mode: A commutative and associative operation is applied (+, \*, min, max, and, or, xor etc.)



# Memory conflicts in PRAM

During the execution of a PRAM algorithm, multiple PUs could try to access to the same memory location (memory conflict).



- Possible resolutions?
- **EREW:** Exclusive Read Exclusive Write
  - Simultaneous read/write of the same memory location by multiple PUs is forbidden.
  - The most restrictive model
- **ERCW:** Exclusive Read, Concurrent Write
  - Does not exist. Why?
    - Does not make sense architecturally.

# A first PRAM algorithm: Search in an array

Given an array  $A[N]$  and a key  $e$ , find the unique element index in  $A$  that contains  $e$ .

---

SEARCHINDEX( $A[N]$ ,  $e$ )

---

```
1:  $idx = 0$ 
2: forall  $i \leftarrow 1 \dots N$  in parallel do
3:   if  $e = A[i]$  then
4:      $idx \leftarrow i$ 
5: return  $idx$ 
```

---

- Sequential algorithm/complexity?
  - Iterate over all array elements,  $O(N)$ .
- PRAM complexity?
  - $O(1)$  time complexity since all comparisons are done simultaneously.
  - Is EREW sufficient? Or do we need CREW, CRCW?
    - EREW OK as long as the element is unique.
    - What about reading the variable  $N$ ?

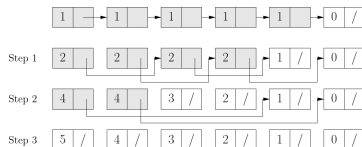
## Another algorithm: List ranking

Given a linked list  $next[N]$ , find the distances from the end of the list for each node.

- $next[i] = \mathbf{Nil} \implies d[i] = 0$
- $next[i] = j \neq \mathbf{Nil} \implies d[i] = d[j] + 1$
- Sequential algorithm? Complexity?
  - Recursion + memoisation in  $O(N)$ .

# List ranking (cont.)

How to parallelize in PRAM?



- Idea: Recursive pointer jumping
  - First, each pointer jump adds 1 distance
  - Then, as we jump over pointers, we add the distance of that pointer.
  - Continue until all pointers point to the end of list.
  - How many iterations until convergence?
    - At each jump, either jump distance of a pointer doubles, or it points to the end of list.
    - In  $\log_2 N$  iterations, all pointers point to the end of list, with correct jump distance  $d[i]$ .

## List ranking (cont.)

## How to write the PRAM algorithm?

```

1 RANK_COMPUTATION( $L$ )
2   forall  $i$  in parallel do                                     { Initialization }
3     if  $next[i] = Nil$  then  $d[i] \leftarrow 0$  else  $d[i] \leftarrow 1$ 
4   while there exists a node  $i$  such that  $next[i] \neq Nil$  do { Main loop }
5     forall  $i$  in parallel do
6       if  $next[i] \neq Nil$  then
7          $d[i] \leftarrow d[i] + d[next[i]]$ 
8          $next[i] \leftarrow next[next[i]]$ 

```

- Initialization is done in  $O(1)$ .
- Computation terminates in  $\log_2 N$  iterations using  $N$  PUs.



## List ranking (cont.)

Read/write conflicts in lines 7, 8?

```

1  RANK_COMPUTATION( $L$ )
2  forall  $i$  in parallel do                                     { Initialization }
3  |   if  $next[i] = Nil$  then  $d[i] \leftarrow 0$  else  $d[i] \leftarrow 1$ 
4  while there exists a node  $i$  such that  $next[i] \neq Nil$  do { Main loop }
5  |   forall  $i$  in parallel do
6  |   |   if  $next[i] \neq Nil$  then
7  |   |   |    $d[i] \leftarrow d[i] + d[next[i]]$ 
8  |   |   |    $next[i] \leftarrow next[next[i]]$ 

```

- We consider it as a sequence of synchronous operations:
  - $temp1[i] \leftarrow d[next[i]]$
  - $temp2[i] \leftarrow d[i] + temp1[i]$
  - $d[i] \leftarrow temp2[i]$
- Asymptotically the same!
- No read nor write conflicts! EREW is sufficient.

## List ranking (cont.)

## How to know when to stop?

```

1  RANK_COMPUTATION( $L$ )
2  forall  $i$  in parallel do                                     { Initialization }
3  |   if  $next[i] = \text{Nil}$  then  $d[i] \leftarrow 0$  else  $d[i] \leftarrow 1$ 
4  while there exists a node  $i$  such that  $next[i] \neq \text{Nil}$  do { Main loop }
5  |   forall  $i$  in parallel do
6  | |   if  $next[i] \neq \text{Nil}$  then
7  | | |    $d[i] \leftarrow d[i] + d[next[i]]$ 
8  | | |    $next[i] \leftarrow next[next[i]]$ 

```

- CRCW PRAM?
  - Write  $done \leftarrow (next[i] = \text{Nil})$  at the end of while
  - For write conflicts, use either “min” or “and” fusion mode
- CREW PRAM?
  - Each process sets  $done[i] \leftarrow (next[i] = \text{Nil})$  at the end of loop
  - Run the reduction algorithm to compute the global  $done$  ( $O(\log N)$  CREW, à venir).

## List ranking (cont.)

## How to know when to stop?

```

1  RANK_COMPUTATION( $L$ )
2  forall  $i$  in parallel do                                { Initialization }
3      if  $next[i] = Nil$  then  $d[i] \leftarrow 0$  else  $d[i] \leftarrow 1$ 
4  while there exists a node  $i$  such that  $next[i] \neq Nil$  do { Main loop }
5      forall  $i$  in parallel do
6          if  $next[i] \neq Nil$  then
7               $d[i] \leftarrow d[i] + d[next[i]]$ 
8               $next[i] \leftarrow next[next[i]]$ 

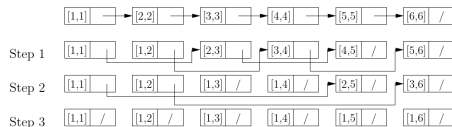
```

- EREW PRAM?

- Each process sets  $done[i] \leftarrow (next[i] = \text{Nil})$  at the end of loop
- Run the reduction algorithm to compute the global  $done$  ( $O(\log N)$  CREW, à venir).
- Run a broadcast algorithm to copy the global  $done$  in each  $done[i]$  ( $O(\log N)$  en EREW, à venir).
- Complexity?  $O(\log N \log N) = O(\log^2 N)$ ?
  - Perform this check every  $\log N$  iterations.
  - Amortized cost per iteration is  $O(1)$ .

# Prefix sum computation on a list

Given a linked list of values  $(x_1, \dots, x_N)$  and a binary associative operation  $\otimes$ , compute the sequence  $(y_1, \dots, y_N)$  such that  $y_k = x_1 \otimes \dots \otimes x_k$



**FIGURE 1.3:** Example execution of the prefix computation algorithm.  $[i, j]$  denotes  $x_i \otimes x_{i+1} \otimes \dots \otimes x_j$  for  $i \leq j$ .

- $x$  is in a linked list form
  - $x[i]$  gives the value
  - $next[i]$  gives the index to the next element
- Use the pointer jumping technique
  - Each element  $i$  applies itself to its  $next[i]$  before jumping
  - $O(\log N)$  iterations maximum

## Prefix sum computation on a list (cont.)

Given a linked list of values  $(x_1, \dots, x_N)$  and a binary associative operation  $\otimes$ , compute the sequence  $(y_1, \dots, y_N)$  such that  $y_k = x_1 \otimes \dots \otimes x_k$

```

1 PREFIX_COMPUTATION( $L$ )
2   forall  $i$  in parallel do                                     { Initialization }
3      $y[i] \leftarrow x[i]$ 
4   while there exists a node  $i$  such that  $next[i] \neq Nil$  do { Main loop }
5     forall  $i$  in parallel do
6       if  $next[i] \neq Nil$  then
7          $y[next[i]] \leftarrow y[i] \otimes y[next[i]]$ 
8          $next[i] \leftarrow next[next[i]]$ 

```

- $x$  is in a linked list form
  - $x[i]$  gives the value
  - $next[i]$  gives the index to the next element
- Use the pointer jumping technique
  - Each element  $i$  applies itself to its  $next[i]$  before jumping
  - $O(\log N)$  iterations maximum
  - What type of PRAM is needed?

## Evaluating the performance of PRAM algorithms

For a problem  $A(n)$  of size  $n$  and corresponding sequential and parallel algorithms, we define the following metrics for evaluation:

- $T_{seq}(n)$ : Execution time of the best sequential algorithm
- $T(n, p)$ : Parallel execution time using  $p$  PUs
- $C(n, p) = pT(n, p)$ : Parallel execution cost using  $p$  PUs
- $W(n, p)$ : Total work performed in parallel (total number of operations across all PUs)

## Evaluating the performance of PRAM algorithms

For a problem  $A(n)$  of size  $n$  and corresponding sequential and parallel algorithms, we define the following metrics for evaluation:


- In general, when algorithm uses the maximum number of PUs for the problem (e.g.,  $p = n$ ), for simplicity, we just use  $W(n) = W(n, p)$  (same for  $T(n)$  and  $C(n)$ )
- $D(n) = \lim_{p \rightarrow \infty} T(n, p)$ : Depth of the parallel algorithm
- $T_{seq}(n) \leq W(n, p) \leq C(n, p)$ : Parallelism potentially gives an overhead
- $S(n, p) = \frac{T_{seq}(n)}{T(n, p)}$ : Speedup
- $E(n, p) = \frac{S(n, p)}{p} = \frac{T_{seq}(n)}{pT(n, p)} = \frac{T_{seq}(n)}{C(n, p)}$ : Efficiency
- If  $W(n, p) = O(T_{seq}(n))$ , algorithm is called **work-optimal**.
- If  $C(n, p) = O(T_{seq}(n))$ , algorithm is called **efficient**.

## A simple simulation

Let  $A$  be an algorithm whose execution time is  $t$  using PRAM with  $p$  PUs. Then,  $A$  can be simulated with a same type PRAM with  $p' \leq p$  PUs in  $O(\frac{tp}{p'})$  time. The cost of the algorithm on this latter PRAM (having  $p'$  PUs) is at most 2x the cost of the initial PRAM with  $p$  PUs.

- Each "big" step of the PRAM can be simulated in  $\lceil \frac{p}{p'} \rceil$  substeps.
- There are at most  $t$  such substeps, thus  $t' = O(\frac{p}{p'} t) = O(\frac{tp}{p'})$ .
- New cost :

$$C_{p'} = t'p' \leq \lceil \frac{p}{p'} \rceil p't \leq (\frac{p}{p'} + 1)p't = pt(1 + \frac{1}{p'}) = C_p(1 + 1/p') \leq 2C_p$$

- **Corollary:** The cost of a parallel execution of a PRAM algorithm is asymptotically the same or superior to the best sequential algorithm
  - If the cost of PRAM algorithm is asymptotically equivalent to its sequential time, then the algorithm is called **efficient**.
- 



# Brent theorem

Let  $A$  be an algorithm that performs  $m$  operations in  $t$  time using PRAM (with a certain, potentially infinite, number of PUs). Then,  $A$  can be simulated with a PRAM of the same type having  $p$  PUs in time  $O(\frac{m}{p} + t)$ .

- Let  $m_i$  be the number of operations performed in the step  $i$  ( $\sum_{i=1}^t m_i = m$ ).
- Each coarse step of PRAM having  $m_i$  operations could be simulated in  $\lceil \frac{m_i}{p} \rceil$  time using a PRAM with  $p$  PUs.

•

$$t' = \sum_{i=1}^t \lceil \frac{m_i}{p} \rceil \leq \sum_{i=1}^t (\frac{m_i}{p} + 1) = \frac{m}{p} + t$$

## Comparison between the power of PRAM models

Is there a problem that one can solve asymptotically faster using CRCW PRAM instead of CREW PRAM

```

1  COMPUTE_MAXIMUM( $A, n$ )
2    forall  $i \in \{1, \dots, n\}$  in parallel do
3       $m[i] \leftarrow \text{True}$ 
4    forall  $i, j \in \{1, \dots, n\}^2, i \neq j$  in parallel do
5      if  $A[i] < A[j]$  then  $m[i] \leftarrow \text{False}$ 
6    forall  $i \in \{1, \dots, n\}$  in parallel do
7      if  $m[i] = \text{True}$  then  $max \leftarrow A[i]$ 
8    return  $max$ 

```

- Find the maximum of an array  $A[N]$
- Can be performed in  $O(1)$  time using CRCW
- CREW requires at least  $O(\log_2 n)$  due to reduction

## Comparison between the power of PRAM models (cont.)

Is there a problem that one can solve asymptotically faster using CREW PRAM instead of EREW PRAM

---

SEARCH(A[n], e)

```
1:  $idx = 0$ 
```

2: **forall**  $i = 1, \dots, N$  **in parallel do**

3:   **if**  $e = A[i]$  **then**

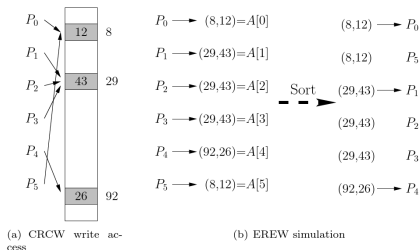
4:  $idx \leftarrow i$ 

```
5: return idx
```

- Search an element  $e$  in an array (with unique elements)  $A[N]$
- Could be performed in  $O(1)$  using CREW
- EREW requires at least  $O(\log_2 n)$  for diffusing  $e$  to all PUs.

# Simulation theorem

The execution time of a CRCW PRAM algorithm using  $P$  PUs is at most  $O(\log N)$  times faster than the best EREW PRAM algorithm using  $P$  PUs for the same problem



- To simulate each step of CRCW using EREW in consistent mode
- For each write at an index, put (index, value) pair in an array  $A[P]$ .
- Sort  $A[P]$  ( $O(\log P)$  in EREW PRAM)
- Perform the write for  $A[0]$ , then for  $A[i]$  if  $A[i] \neq A[i - 1]$ .



## **Contact**

Oguz Kaya

Université Paris-Saclay and LRI, Paris, France

[oguz.kaya@lri.com](mailto:oguz.kaya@lri.com)

[www.oguzkaya.com](http://www.oguzkaya.com)